



**University of Duisburg-Essen**  
Networked Embedded Systems Group  
Institute for Computer Science and  
Business Information Systems (ICB)  
Schützenbahn 70  
D-45127 Essen, Germany

# **Kommunikationsnetze 2**

## **7 – Infrastructure**

Prof. Dr. Pedro José Marrón

# IP Addresses: How to Get One?

---


- How does a host get an IP address?
- Hard-coded by system admin in a file
- DHCP: Dynamic Host Configuration Protocol
  - Dynamically get address from a server
    - “plug-and-play”

# DHCP: Dynamic Host Configuration Protocol

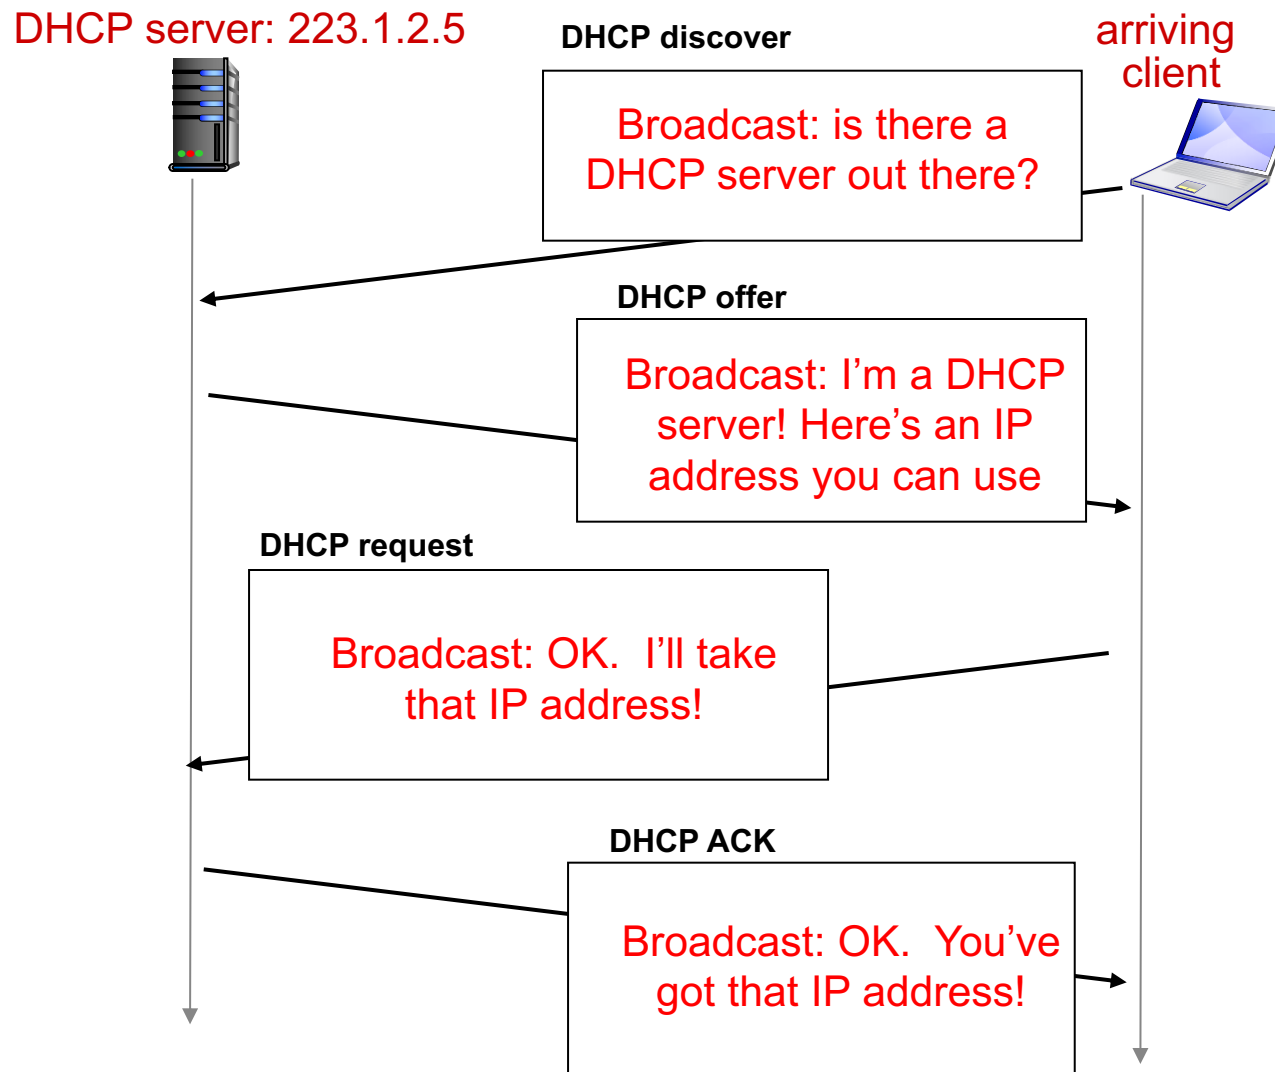
---

- Goal: allow host to dynamically obtain its IP address from network server when it joins network
  - Can renew its lease on address in use
  - Allows reuse of addresses (only hold address while connected/“on”)
  - Support for mobile users who want to join network (more shortly)
- DHCP overview
  - Host broadcasts “DHCP discover” msg [optional]
  - DHCP server responds with “DHCP offer” msg [optional]
  - Host requests IP address: “DHCP request” msg
  - DHCP server sends address: “DHCP ack” msg

# DHCP Client-Server Scenario



# DHCP Client-Server Scenario




# DHCP: More than IP Addresses

---


- DHCP can return more than just allocated IP address on subnet
  - Address of first-hop router for client
  - Name and IP address of DNS server
  - Network mask (indicating network versus host portion of address)

# DHCP: Example



- Connecting laptop needs its IP address, addr of first-hop router, addr of DNS server: use DHCP
- DHCP request encapsulated in UDP, encapsulated in IP, encapsulated in 802.1 Ethernet
- Ethernet frame broadcast (dest: FFFFFFFFFFFFFFFF) on LAN, received at router running DHCP server
- Ethernet demuxed to IP, demuxed to UDP, demuxed to DHCP

# DHCP: Example



- DHCP server formulates DHCP ACK containing client's IP address, IP address of first-hop router for client, name & IP address of DNS server
- Encapsulation of DHCP server, frame forwarded to client, demuxing up to DHCP at client
- Client now knows its IP address, name and IP address of DNS server, IP address of its first-hop router



# DHCP: Wireshark Output (Home LAN)

Message type: **Boot Request (1)**

Hardware type: Ethernet

Hardware address length: 6

Hops: 0

**Transaction ID: 0x6b3a11b7**

Seconds elapsed: 0

Bootp flags: 0x0000 (Unicast)

Client IP address: 0.0.0.0 (0.0.0.0)

Your (client) IP address: 0.0.0.0 (0.0.0.0)

Next server IP address: 0.0.0.0 (0.0.0.0)

Relay agent IP address: 0.0.0.0 (0.0.0.0)

**Client MAC address: Wistron\_23:68:8a (00:16:d3:23:68:8a)**

Server host name not given

Boot file name not given

Magic cookie: (OK)

Option: (t=53,l=1) **DHCP Message Type = DHCP Request**

Option: (61) Client identifier

Length: 7; Value: 010016D323688A;

Hardware type: Ethernet

Client MAC address: Wistron\_23:68:8a (00:16:d3:23:68:8a)

Option: (t=50,l=4) Requested IP Address = 192.168.1.101

Option: (t=12,l=5) Host Name = "nomad"

**Option: (55) Parameter Request List**

Length: 11; Value: 010F03062C2E2F1F21F92B

**1 = Subnet Mask; 15 = Domain Name**

**3 = Router; 6 = Domain Name Server**

**44 = NetBIOS over TCP/IP Name Server**

.....

request

Message type: **Boot Reply (2)**

Hardware type: Ethernet

Hardware address length: 6

Hops: 0

**Transaction ID: 0x6b3a11b7**

Seconds elapsed: 0

Bootp flags: 0x0000 (Unicast)

**Client IP address: 192.168.1.101 (192.168.1.101)**

Your (client) IP address: 0.0.0.0 (0.0.0.0)

**Next server IP address: 192.168.1.1 (192.168.1.1)**

Relay agent IP address: 0.0.0.0 (0.0.0.0)

Client MAC address: Wistron\_23:68:8a (00:16:d3:23:68:8a)

Server host name not given

Boot file name not given

Magic cookie: (OK)

**Option: (t=53,l=1) DHCP Message Type = DHCP ACK**

**Option: (t=54,l=4) Server Identifier = 192.168.1.1**

**Option: (t=1,l=4) Subnet Mask = 255.255.255.0**

**Option: (t=3,l=4) Router = 192.168.1.1**

**Option: (6) Domain Name Server**

**Length: 12; Value: 445747E2445749F244574092;**

**IP Address: 68.87.71.226;**

**IP Address: 68.87.73.242;**

**IP Address: 68.87.64.146**

**Option: (t=15,l=20) Domain Name = "hsd1.ma.comcast.net."**

reply

# IP Addresses: How to Get One?

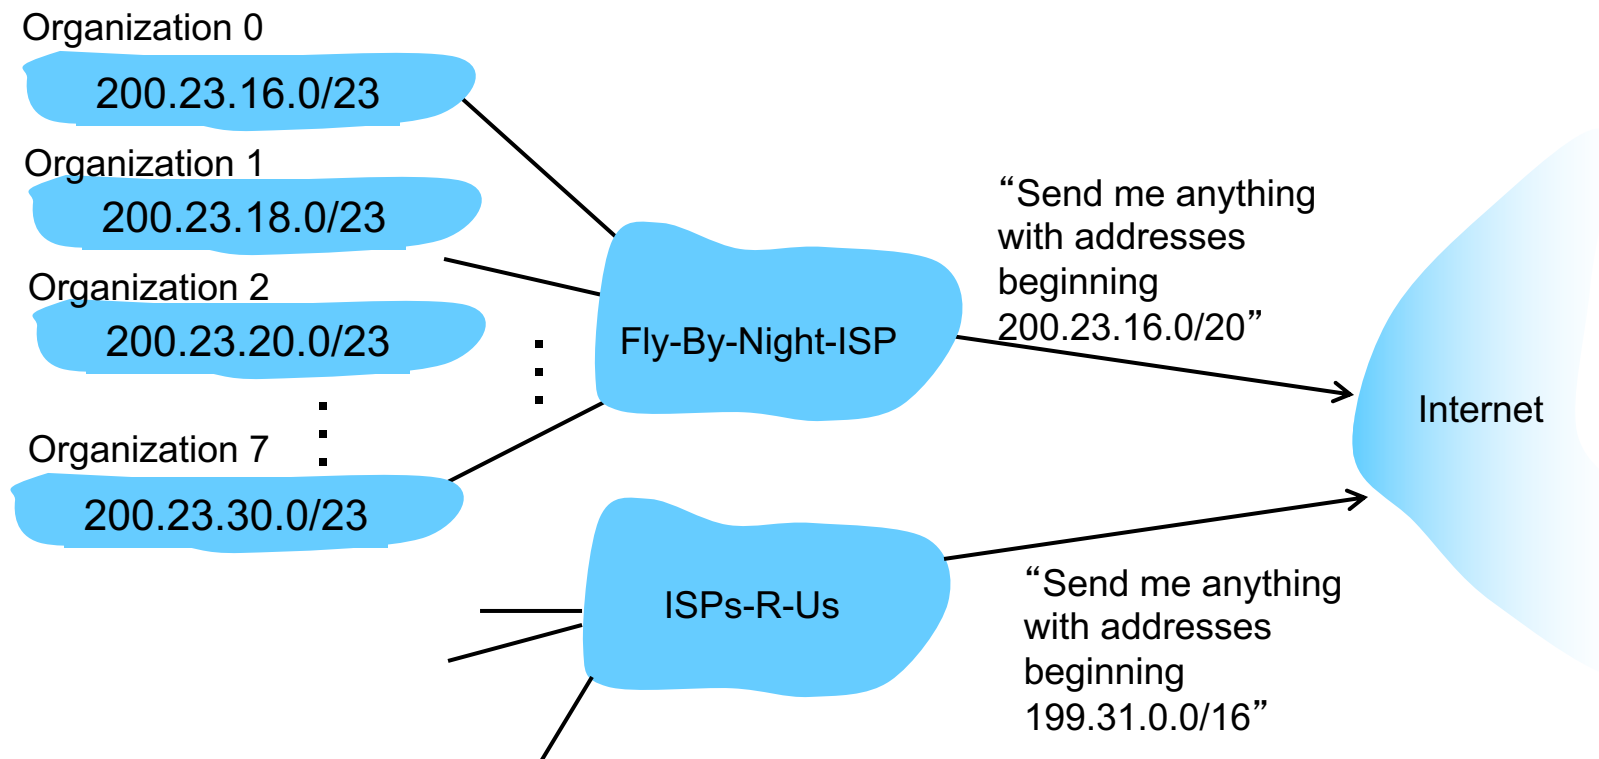
---

- How does network get subnet part of IP addr?
- Gets allocated portion of its provider ISP's address space

ISP's block	<u>11001000</u>	<u>00010111</u>	<u>00010000</u>	00000000	200.23.16.0/20
Organization 0	<u>11001000</u>	<u>00010111</u>	<u>00010000</u>	00000000	200.23.16.0/23
Organization 1	<u>11001000</u>	<u>00010111</u>	<u>00010010</u>	00000000	200.23.18.0/23
Organization 2	<u>11001000</u>	<u>00010111</u>	<u>00010100</u>	00000000	200.23.20.0/23
...		.....		....	....
Organization 7	<u>11001000</u>	<u>00010111</u>	<u>00011110</u>	00000000	200.23.30.0/23

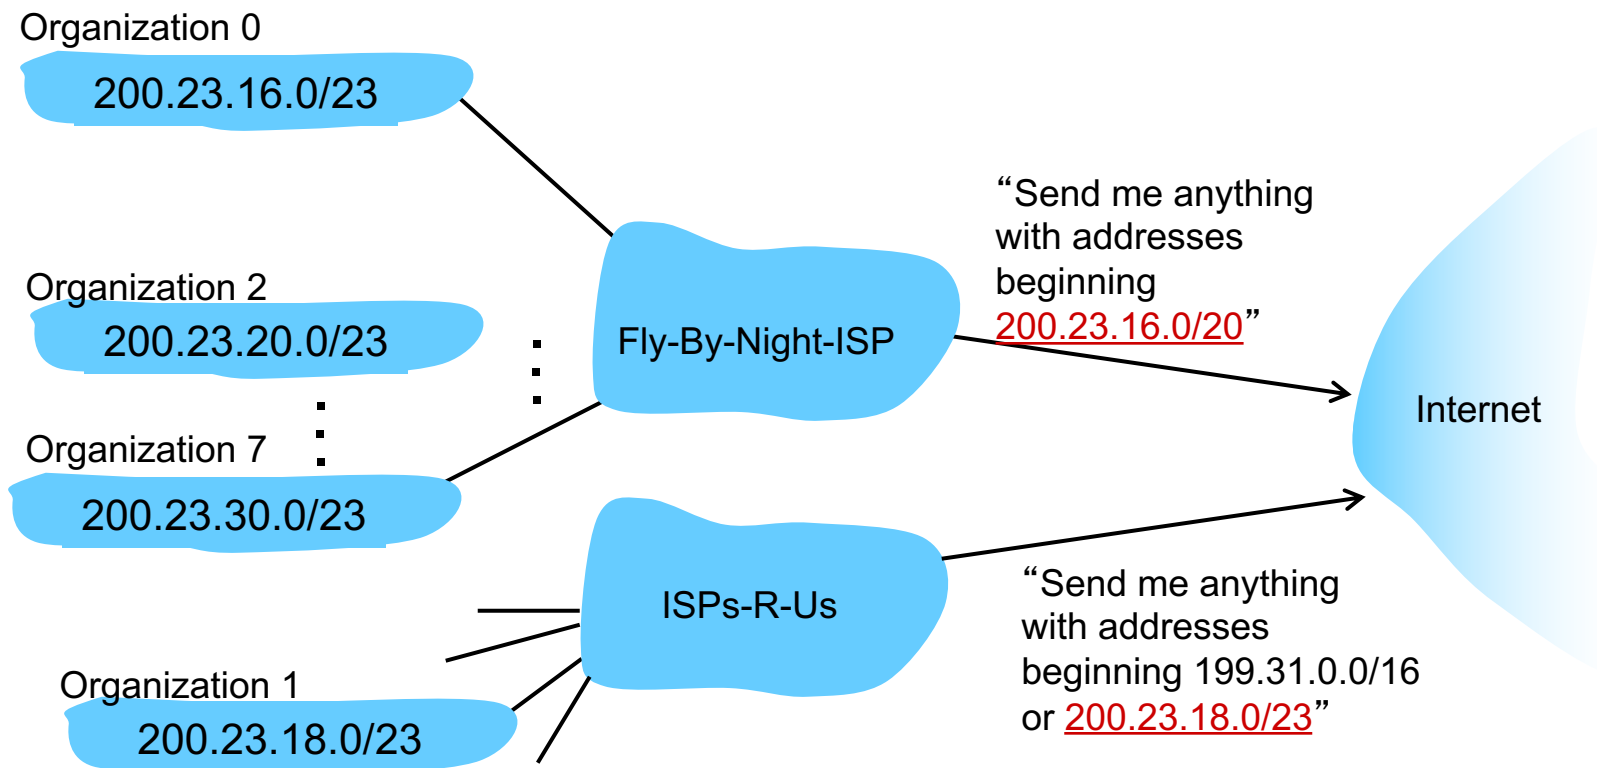
# Hierarchical Addressing: Route Aggregation

- Hierarchical addressing allows efficient advertisement of routing information



# Hierarchical Addressing: More Specific Routes

- ISPs-R-Us has more specific route to Organization 1




# IP Addressing: The Last Word...

---

- How does an ISP get block of addresses?
- ICANN: Internet Corporation for Assigned Names and Numbers
  - <http://www.icann.org>
  - Allocates addresses
  - Manages DNS
  - Assigns domain names, resolves disputes

# NAT/PAT: Network Address Translation



*all* datagrams *leaving* local network have *same* single source NAT IP address: 138.76.29.7, different source port numbers

datagrams with source or destination in this network have 10.0.0/24 address for source, destination (as usual)

# NAT/PAT: Network Address Translation

---

- Motivation: Local network uses just one IP address as far as outside world is concerned
  - Range of addresses not needed from ISP: just one IP address for all devices
  - Can change addresses of devices in local network without notifying outside world
  - Can change ISP without changing addresses of devices in local network
  - Devices inside local net not explicitly addressable, visible by outside world (a security bonus)

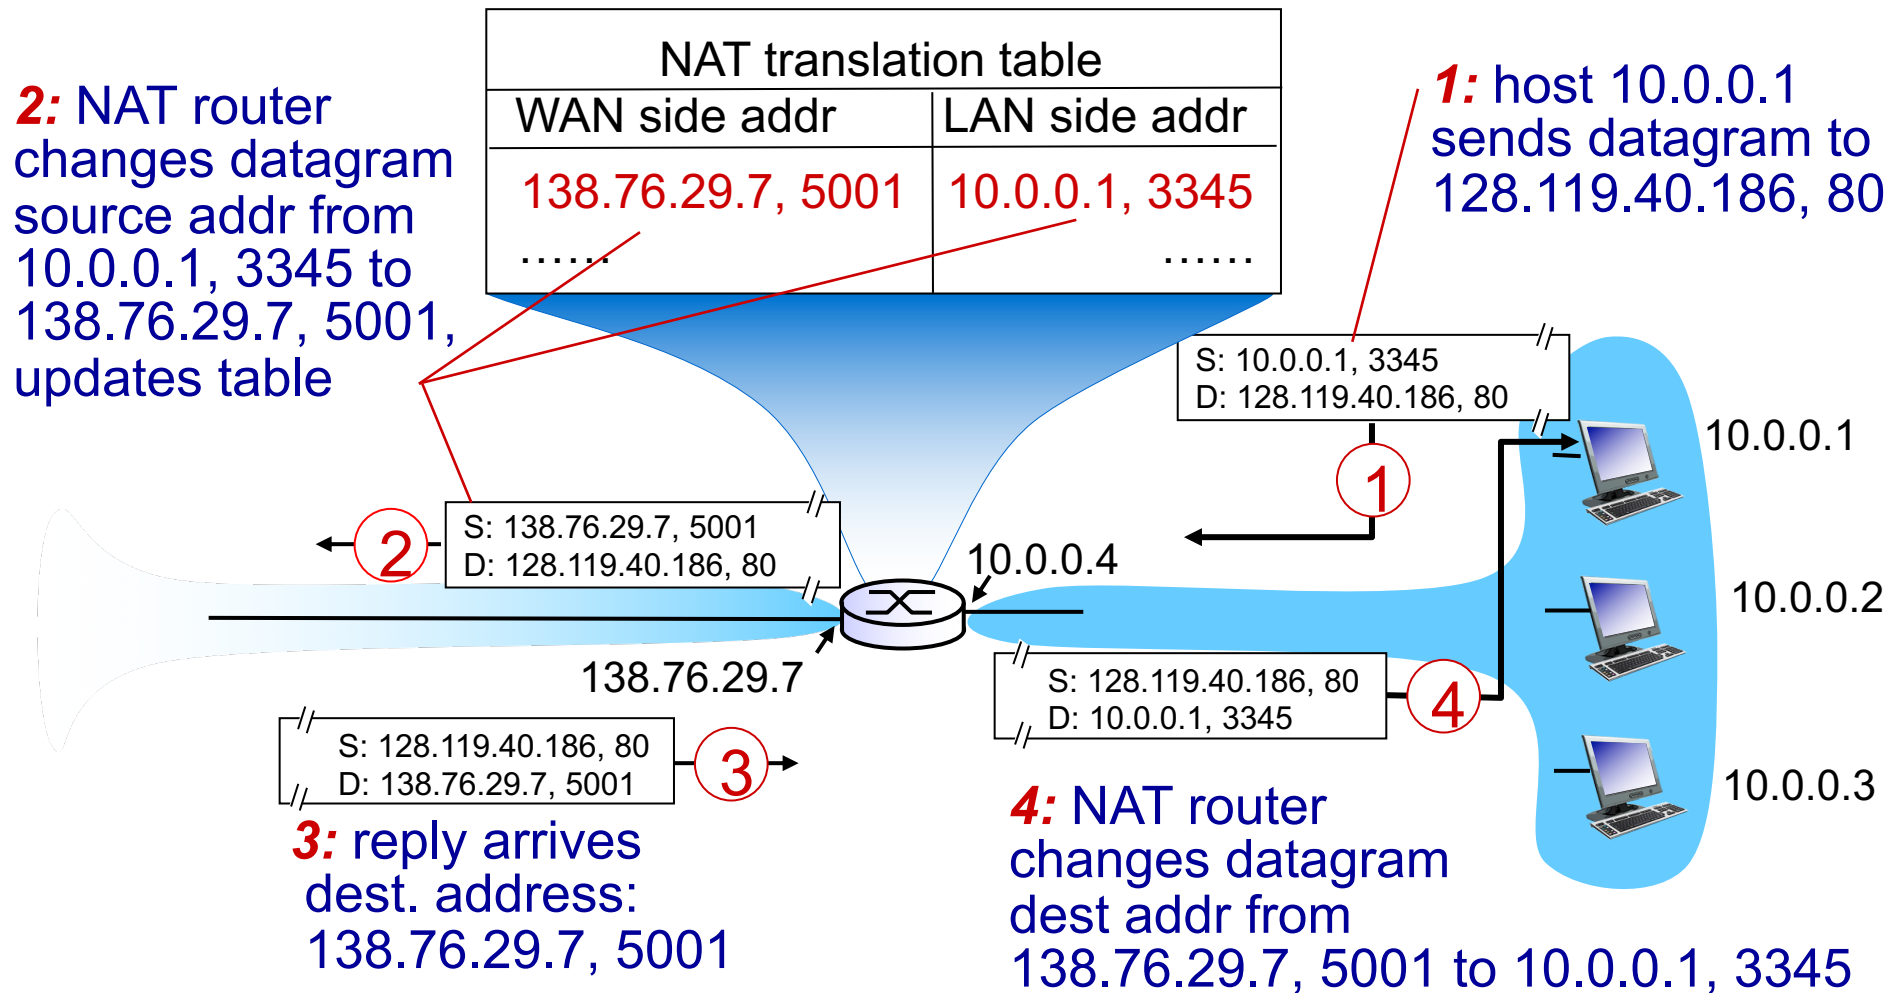
# NAT/PAT: Implementation

---

- Implementation: NAT router must
  - Replace (source IP address, port #) of every outgoing datagram to (NAT IP address, new port #)
    - Remote clients/servers will respond using (NAT IP address, new port #) as destination addr
  - Remember (in NAT translation table) every (source IP address, port #) to (NAT IP address, new port #) translation pair
  - Replace (NAT IP address, new port #) in dest fields of every incoming datagram with corresponding (source IP address, port #) stored in NAT table



# NAT/PAT: Example



# NAT/PAT: Network Address Translation

---

- 16-bit port-number field
  - 60.000 simultaneous connections with a single LAN-side address
- NAT is controversial
  - Routers should only process up to layer 3
  - Address shortage should be solved by IPv6
  - Violates end-to-end argument
    - NAT possibility must be taken into account by app designers, e.g., P2P applications
  - NAT traversal: what if client wants to connect to server behind NAT?

# NAT: Types

---

- PAT (or NAT Overload)
  - Mapping between multiple local IP addresses to a single public IP address (many-to-one) using different Layer 4 ports
  - (What we have seen until now)
- Static NAT
  - One-to-one mapping between a local IP address and a global IP address
  - Static configuration allowing services to be accessed from outside
- Dynamic NAT
  - A pool of public IP addresses is available
  - At run-time, local IP addresses are mapped to one of the free public IP addresses in the pool
  - After a timeout, the mapping is purged from the translation table

# DNS: Domain Name System

---

- Internet hosts, routers identifiers:
  - IP address – used for addressing datagrams
  - “name”, e.g., uni-due.de – used by humans
- How to map between IP address and name, and vice versa?
- Domain Name System
  - Distributed database implemented in hierarchy of many name servers
  - Application-layer protocol: hosts, name servers communicate to resolve names (address/name translation)
    - Note: core Internet function, implemented as application-layer protocol


# DNS: Services, Structure

---

- DNS services
  - Hostname to IP address translation
  - Host aliasing
    - Canonical, alias names
  - Mail server aliasing
  - Load distribution
    - Replicated Web servers
    - Many IP addresses correspond to one name
- Why not centralize DNS?
  - Single point of failure
  - Traffic volume
  - Distant centralized database
  - Maintenance
  - Does not scale!

# DNS: A Distributed, Hierarchical Database

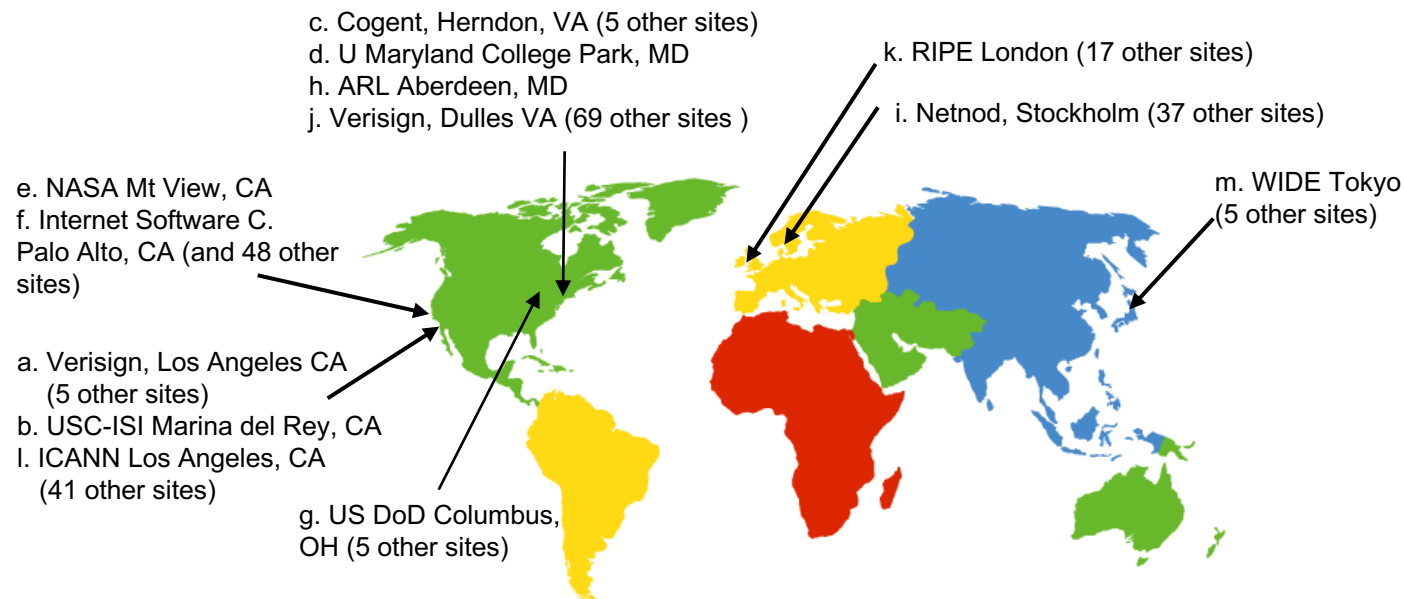
---



- Client wants IP for nes.uni-due.de; 1<sup>st</sup> approximation:
  - Client queries root server to find de DNS server
  - Client queries de DNS server to get uni-due.de DNS server
  - Client queries uni-due.de DNS server to get IP address of nes.uni-due.de

# DNS: Root Name Servers

- Contacted by local name server that can not resolve name
- Root name server:
  - Contacts authoritative name server if name mapping not known
  - Gets mapping
  - Returns mapping to local name server
- 13 logical root name “servers” worldwide
  - Each server replicated multiple times



# Top-Level Domain and Authoritative Servers

---

- Top-level domain (TLD) servers
  - Responsible for com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g., uk, de, fr, ca, it, jp
  - Verisign (after Network Solutions) maintains servers for .com TLD
  - DENIC eG for .de TLD
- Authoritative DNS servers
  - Organization's own DNS server(s) providing authoritative hostname to IP mappings for organization's named hosts
  - Can be maintained by organization or service provider



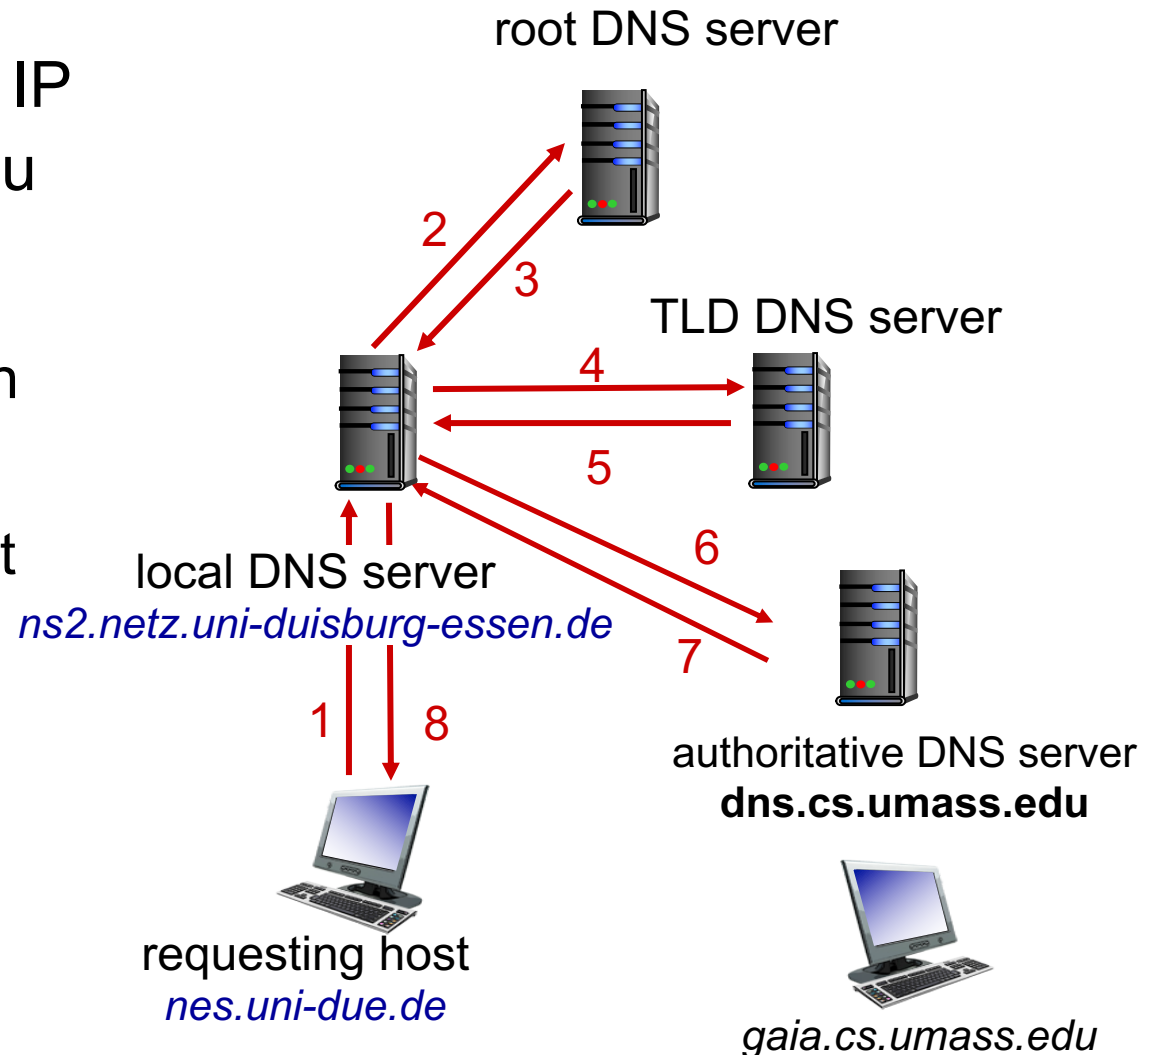
# Local DNS Name Server

---

- Does not strictly belong to hierarchy
- Each ISP (residential ISP, company, university) has one
  - Also called “default name server”
- When host makes DNS query, query is sent to its local DNS server
  - Has local cache of recent name-to-address translation pairs (but may be out-of-date)
  - Acts as proxy, forwards query into hierarchy

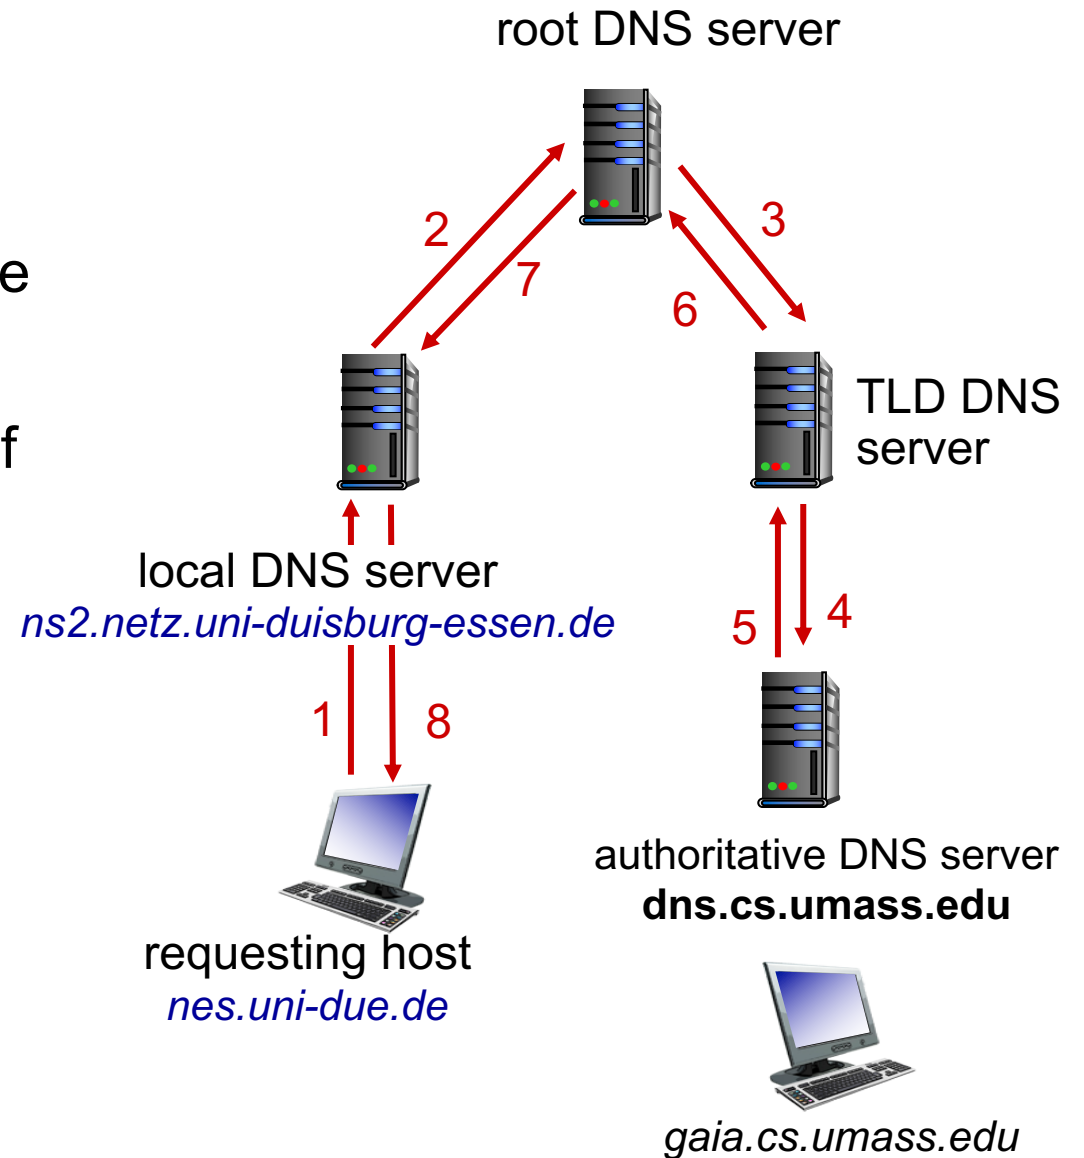
# DNS Name Resolution Example

- Host at nes.uni-due.de wants IP address for gaia.cs.umass.edu
- Iterated query:
  - Contacted server replies with name of server to contact
  - “I do not know this name, but ask this server”



# DNS Name Resolution Example

- Recursive query:
  - Puts burden of name resolution on contacted name server
  - Heavy load at upper levels of hierarchy?



# DNS Lookup Trace Example

```
dig +nocmd +additional +trace +nodnssec www.nes.uni-due.de A
```

```
.           364237      IN          NS          e.root-servers.net.
.           364237      IN          NS          a.root-servers.net.
.           364237      IN          NS          d.root-servers.net.
.           364237      IN          NS          c.root-servers.net.
.           364237      IN          NS          b.root-servers.net.
.           364237      IN          NS          l.root-servers.net.
.           364237      IN          NS          f.root-servers.net.
.           364237      IN          NS          i.root-servers.net.
.           364237      IN          NS          g.root-servers.net.
.           364237      IN          NS          m.root-servers.net.
.           364237      IN          NS          j.root-servers.net.
.           364237      IN          NS          h.root-servers.net.
.           364237      IN          NS          k.root-servers.net.
;; Received 239 bytes from 134.91.76.7#53(134.91.76.7) in 1 ms

de.         172800      IN          NS          a.nic.de.
de.         172800      IN          NS          f.nic.de.
de.         172800      IN          NS          l.de.net.
de.         172800      IN          NS          n.de.net.
de.         172800      IN          NS          s.de.net.
de.         172800      IN          NS          z.nic.de.
;; Received 389 bytes from 192.58.128.30#53(j.root-servers.net) in 15 ms

uni-due.de. 86400         IN          NS          dns-2.dfn.de.
uni-due.de. 86400         IN          NS          ns1.uni-duisburg-essen.de.
uni-due.de. 86400         IN          NS          ns2.uni-duisburg-essen.de.
;; Received 258 bytes from 194.246.96.1#53(z.nic.de) in 6 ms

nes.uni-due.de. 86400      IN          NS          ns1.netz.uni-duisburg-essen.de.
nes.uni-due.de. 86400      IN          NS          ns2.netz.uni-duisburg-essen.de.
;; Received 107 bytes from 193.174.75.54#53(dns-2.dfn.de) in 14 ms

www.nes.uni-due.de. 86400      IN          A           134.91.76.7
;; Received 63 bytes from 132.252.1.7#53(ns2.netz.uni-duisburg-essen.de) in 1 ms
```

# DNS: Caching, Updating Records

---

- Once (any) name server learns mapping, it caches mapping
  - Cache entries timeout (disappear) after some time (TTL)
  - TLD servers typically cached in local name servers
    - Thus root name server not often visited
- Cached entries may be out-of-date (best effort name-to-address translation!)
  - If name host changes IP address, may not be known Internet-wide until all TTLs expire
- Update/notify mechanisms proposed IETF standard
  - RFC 2136

# DNS Records

---

- DNS: distributed database storing resource records (RR)

RR format: (**name**, **ttl**, **class**, **type**, **value**)

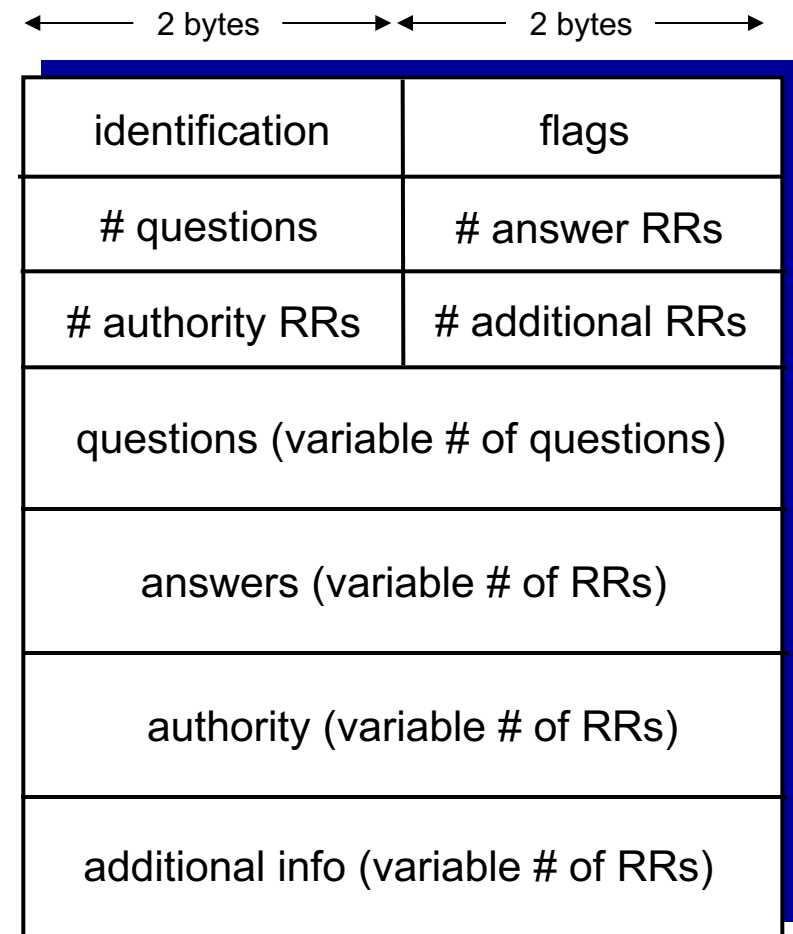
- Type = A
  - **name** is hostname
  - **value** is IP address
- Type = NS
  - **name** is domain
  - **value** is hostname of authoritative name server for this domain
- Type = CNAME
  - **name** is alias name for some “canonical” (the real) name
  - **value** is canonical name
- Type = MX
  - **value** is name of mailserver associated with **name**

# DNS Protocol, Messages

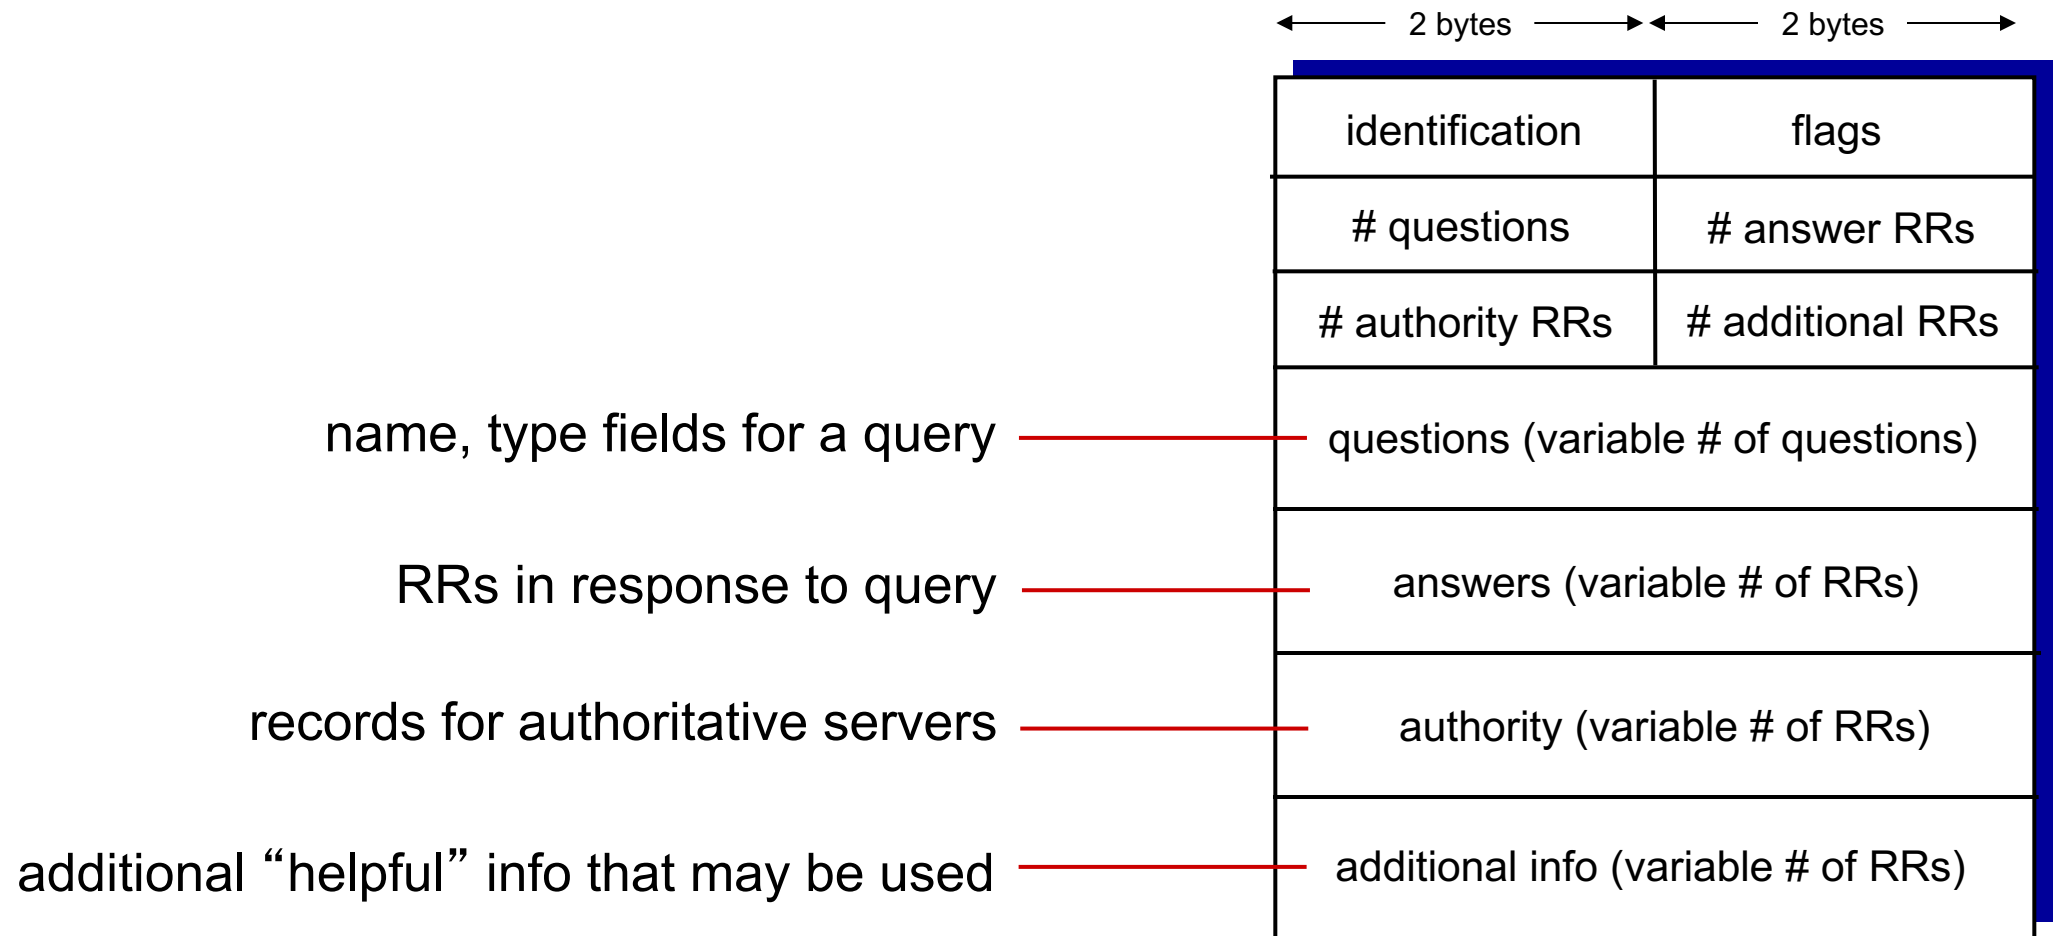
- Query and reply messages, both with same message format

- Message header

- Identification: 16 bit # for query
  - Reply uses the same #
- Flags:
  - Query or reply
  - Recursion desired
  - Recursion available
  - Reply is authoritative



# DNS Protocol, Messages





# Attacking DNS

---

- DDoS attacks
  - Target root servers with traffic
    - Not successful to date
    - Traffic filtering
    - Local DNS servers cache IPs of TLD servers, allowing root server bypass
  - Target TLD servers
    - Potentially more dangerous
- Redirect attacks
  - Main-in-the-middle
    - Intercept queries
  - DNS poisoning
    - Send bogus replies to DNS server, which caches
- Exploit DNS for DDoS
  - Send queries with spoofed source address: target IP
  - Requires amplification